



# Virtual Memory(2)

Dept. of Computer Science

Hanyang University





## Allocation of Frames

- How to allocate page frames to multiple processes?
- Each process needs minimum number of pages
  - HW 측면: IBM 370 – 6 pages to handle SS MOVE instruction:
    - instruction is 6 bytes, might span 2 pages
    - 2 pages to handle **from**
    - 2 pages to handle **to**
  - SW 측면:
    - Loop 내의 page 는 한꺼번에 allocate 되는 것이 유리함
    - 그렇지 않으면 매 loop 마다 page fault – CPU/disk load 심한 불균형
- Two major allocation schemes
  - fixed allocation
  - priority allocation



## Fixed Allocation

- Equal allocation
  - Allocate the same number of frames to each process
  - e.g., (100 frames, 5 processes) 20 pages each
- Proportional allocation
  - Allocate according to the size of process

- $s_i$  = size of process  $p_i$
- $S = \sum s_i$
- $m$  = total number of frames
- $a_i$  = allocation for  $p_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$



## Priority Allocation

- Use a proportional allocation scheme using priorities rather than size
- High priority process
  - Give more memory so that it can finish early
  - (decreased I/O → decreased time in 'waiting' state → early finish)
- If process  $P_i$  generates a page fault,
  - select a victim from one of its frames
  - select a victim from lower priority processes' frames

## Global vs. Local Replacement

- Global replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another.
- Local replacement – each process selects from only its own set of allocated frames.




- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
  - Low CPU utilization
  - Operating system thinks that it needs to increase the degree of multiprogramming
  - Another process added to the system (higher MPD)
- Thrashing
  - A process is busy swapping pages in and out
  - CPU is idle most of the time – low throughput

main()

{ for (l=1, 100000) { A = B + X } }

A
main()
X
B

# Thrashing Diagram




- Why does paging work?  
Locality model
  - Process migrates from one locality to another
  - Localities may overlap
- Why does thrashing occur?  
 $(\sum \text{size of locality}) > (\text{total allocated memory size})$



- 프로그램의 메모리 참조는 고도의 지역성을 가짐
- 임의 시간  $\Delta t$  내에 프로그램의 일부분만을 집중적으로 참조
  - 시간 지역성 (Temporal Locality) : 현재 참조된 메모리가 가까운 미래에도 참조될 가능성이 높음  
ex) loop, subroutine, stack
  - 공간 지역성 (Spatial Locality) : 하나의 메모리가 참조되면 주변의 메모리가 계속 참조될 가능성이 높음  
ex) Array Traversal, 명령의 순차 실행

# Locality In a Memory-Reference Pattern





- 12312312324802480248024803366666633666666666633666666666633666

- 9


# Working-Set model

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$




## Working-Set Model

- $WS(t_i) = \{ \text{pages referenced in } [t_i, t_i - \Delta] \}$
- 만일 페이지 P 가  $t_i$ 에  $WS(t_i)$ 에 속하였으면 keep in memory  
안 속하였으면 out of memory
- 이 원칙에 따라 replace, allocate 를 결정
- 따라서 working set model은 allocate/replace 를 같이 결정함
- 시간에 따라 allocation size 가 달라질 수 있음
- $WS(t_i)$  가 모두 보장되어야만 run, 아니면 suspend


## Keeping Track of the Working Set

- 구현: 매 ref 마다 각 page 들의 최근 reference time 을  $\Delta$  와 비교?  
 → too expensive (space for ref-time field + time for comparison)

- Approximate with interval timer + a reference bit
- Example:  $\Delta = 10K$ ,
  - Timer interrupts: every 5K time units.
  - Keep in memory 2 bits for each page.
  - timer interrupts → copy and resets all reference bits
  - If one of the bits in memory = 1  $\Rightarrow$  page belongs to the working set
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units
- Q: How do you decide window size?



# Page-Fault Frequency Scheme



- Establish “acceptable” page-fault rate
  - If actual rate too low, process loses frame
  - If actual rate too high, process gains frame
- Working set 기법은 page 참조 시마다 페이지 집합을 수정
- PFF 기법은 page fault 발생 시에만 페이지 집합을 수정



## Other Benefits of VM: Process Creation

---

- Virtual memory allows other benefits during process creation:
  - Copy-on-Write
  - Memory-Mapped Files




## Copy-on-Write

---

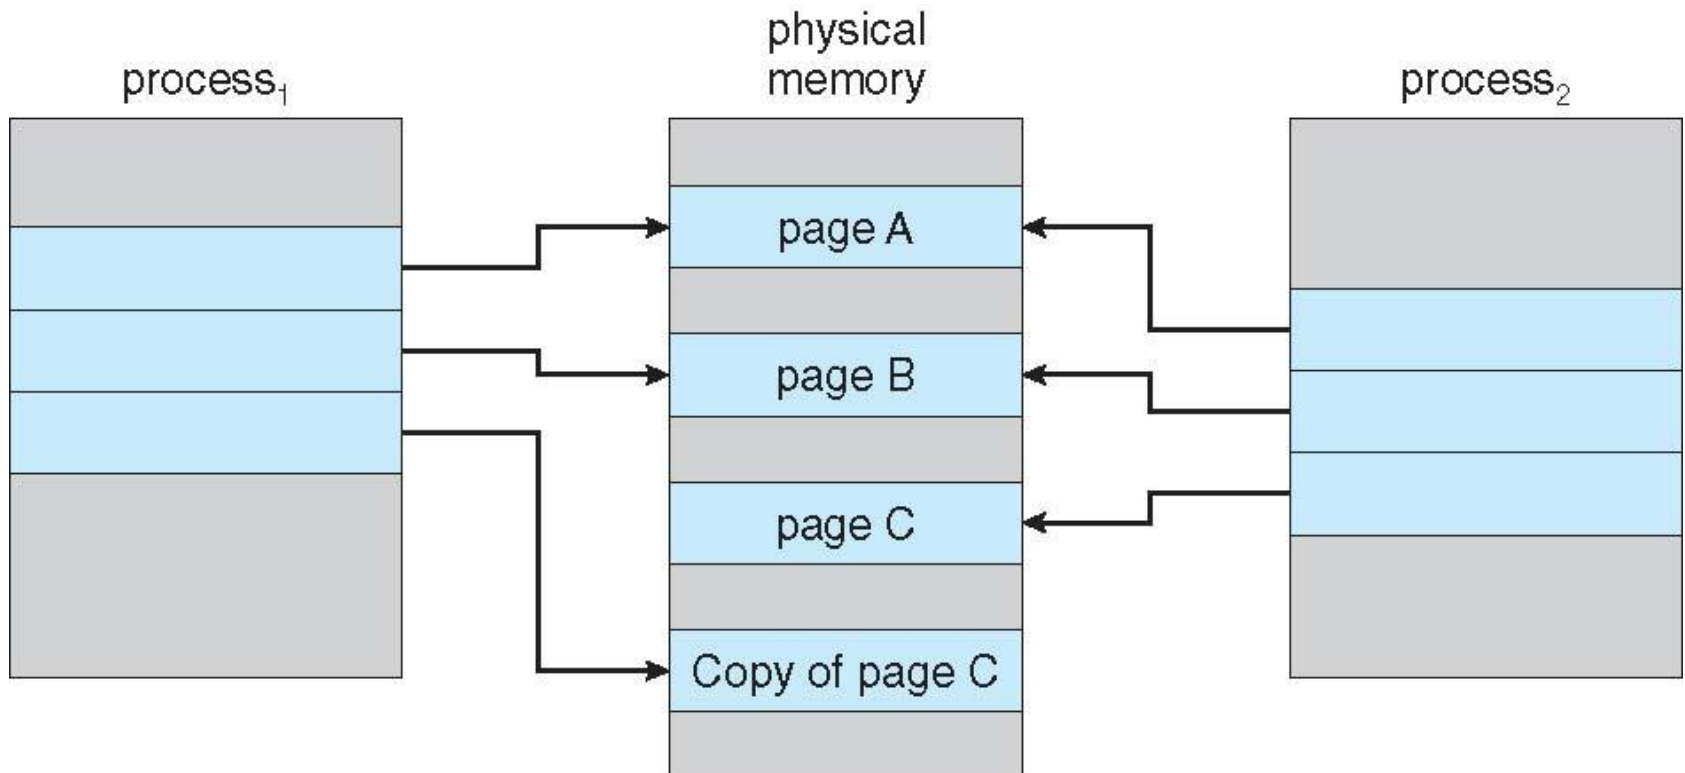
- Copy-on-Write (COW)
  - Allows both parent and child processes to initially *share* the same pages in memory
  - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied

## Before Process 1 Modifies Page C





## After Process 1 Modifies Page C







## Memory-Mapped Files

- Memory-mapped file I/O
  - Allows file I/O to be treated as routine memory access
  - by **mapping** a disk block to a page in memory
- A file is initially read using demand paging
  - A page-sized portion of the file is read from the file system into a physical page
  - Subsequent reads/writes to/from the file are treated as ordinary memory accesses
- Simplifies file accesses
  - Treats file I/O through memory rather than **read()**, **write()** system calls
- Also allows several processes to map the same file allowing the pages in memory to be shared

# Memory Mapped Files





## Other Issues – Prepaging

---

- Prepaging
  - To reduce the large number of page faults that occurs at process startup
  - Prepage all or some of the pages a process will need, before they are referenced
  - But if prepaged pages are unused, I/O and memory was wasted
  - Assume  $s$  pages are prepaged and  $\alpha$  of the pages is used
    - Is cost of  $s * \alpha$  saved pages faults  $>$  or  $<$  than the cost of prepaging  $s * (1 - \alpha)$  unnecessary pages?
    - $\alpha$  near zero  $\Rightarrow$  prepaging loses



## Other Issues – Page size

- Considerations for page size selection
  - Internal fragmentation
  - Page table size
  - Disk transfer efficiency – seek/rotation vs. transfer
  - Frequency of I/O operations
  - Improved Locality
    - Smaller page size isolate only needed info within page
- Trend
  - Larger page size
  - CPU speed, memory capacity – improves faster than disk speed
  - Page fault (relative) penalty is becoming more costly these days



## Other Issues – TLB Reach

---

- TLB Reach
  - The amount of memory accessible from the TLB
  - $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
  - Ideally, the working set of each process is stored in the TLB
    - Otherwise there is a high degree of TLB misses
  - Increase the Page Size
    - This may lead to an increase in fragmentation as not all applications require a large page size
  - Provide Multiple Page Sizes
    - This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation



## Other Issues – Program Structure

- Program structure

- `int data[128][128];`
- Each row is stored in one page
- Assume that the # of free frames for data < 128
- Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i][j] = 0;
```

128 x 128 = 16,384 page faults

- Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i][j] = 0;
```

128 page faults (even when there is only one page frame for data)

## Other Issues – I/O interlock

- **I/O Interlock** – Pages must sometimes be locked into memory
- Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm

