

Locking

Operating system

Introduction

- Race condition
- What is lock?
- Lock implementation
 - Atomic Operation
 - Deadlock
 - Interrupt handler
 - Memory ordering
- Lock usage in xv6
 - usual/unusual case
 - In uniprocessor system

Start from race condition

- Any code that accesses shared data concurrently from multiple CPUs is likely to yield incorrect results or a broken data structure.

Parts of the program where the shared resource is accessed is protected. This protected section is the **critical section** or **critical region**.

```
43 // Increment ref count for file f.
44 struct file*
45 filedup(struct file *f)
46 {
47     if(f->ref < 1)
48         panic("filedup");
49     f->ref++;
50     return f;
51 }
```

filedup function in file.c

Start from race condition

Thread 1	Thread 2		Integer value
			0
read value		←	0
increase value			0
write back		→	1
	read value	←	1
	increase value		1
	write back	→	2

Correct result

Thread 1	Thread 2		Integer value
			0
read value		←	0
	read value	←	0
increase value			0
	increase value		0
write back		→	1
	write back	→	1

Incorrect result

What is Lock?

- a **lock** is a synchronization mechanism for enforcing limits on access to a resource in an environment where there are many threads of execution.
- Locks ensure mutual exclusion so that only one CPU can execute the code at a time.
- So locks can be used to protect critical section and to prevent race condition.

What is Lock? (Cont'd)

Only one CPU can execute the code at a time

```
43 // Increment ref count for file f.
44 struct file*
45 filedup(struct file *f)
46 {
47     acquire(&ftable.lock);
48     if(f->ref < 1)
49         panic("filedup");
50     f->ref++;
51     release(&ftable.lock);
52     return f;
53 }
```

filedup function in file.c

Simple lock code. is it working?

No, We use shared data(lk->locked)
so it can cause race condition!

So, We should execute this
instructions(line 31, 32, 55)
atomically.

```
24 void
25 acquire(struct spinlock *lk)
26 {
27     if(holding(lk))
28         panic("acquire");
29
30     for(;;){
31         if (!lk->locked) {
32             lk->locked = 1;
33             break;
34         }
35     }
36     // Record info about lock acquisition for debugging.
37     lk->cpu = cpu;
38     getcallerpcs(&lk, lk->pcs);
39 }
```

```
46 void
47 release(struct spinlock *lk)
48 {
49     if(!holding(lk))
50         panic("release");
51
52     lk->pcs[0] = 0;
53     lk->cpu = 0;
54
55     lk->locked = 0;
56
57 }
```

Simple lock code. is it working? (Cont'd)

Not likely. Because we didn't consider some issues.

```
24 void
25 acquire(struct spinlock *lk)
26 {
27     if(holding(lk))
28         panic("acquire");
29
30
31     // The xchg is atomic.
32     while(xchg(&lk->locked, 1) != 0)
33         ;
34
35     // Record info about lock acquisition for debugging.
36     lk->cpu = cpu;
37     getcallerpcs(&lk, lk->pcs);
38 }
```

```
46 void
47 release(struct spinlock *lk)
48 {
49     if(!holding(lk))
50         panic("release");
51
52     lk->pcs[0] = 0;
53     lk->cpu = 0;
54
55     // equivalent to lk->locked = 0.
56     // This code can't use a C assignment
57     // since it might not be atomic.
58     asm volatile("movl $0, %0" : "+m"
59                  (lk->locked) : );
59 }
```


Deadlock

- Let's say two code paths in xv6 needs locks A and B
- What happened when two CPUs run each code paths concurrently?

```
24 void
25 function1()
26 {
27     acquire(&A.lock);
28     → acquire(&B.lock);
29     // execute critical section code
30     release(&B.lock);
31     release(&A.lock);
32 }
```

I already acquired
lock A and I am
waiting for lock B

```
24 void
25 function2()
26 {
27     acquire(&B.lock);
28     → acquire(&A.lock);
29     // execute critical section code
30     release(&A.lock);
31     release(&B.lock);
32 }
```

I already acquired
lock B and I am
waiting for lock A

Dreadful Deadlock!

- We are stucked at the code line 28, and we cannot progress forever.
- Like this!



Deadlock prevention

- Therefore, all code paths must acquire locks in the same order.

```
1 void
2 function1()
3 {
4   acquire(&A.lock);
5   acquire(&B.lock);
6   // execute critical section code
7   release(&B.lock);
8   release(&A.lock);
9 }
```

```
11 void
12 function2()
13 {
14   acquire(&A.lock);
15   acquire(&B.lock);
16   // execute critical section code
17   release(&B.lock);
18   release(&A.lock);
19 }
```

- Xv6 has few lock-order chains and the longest chains are only two deep.
- Moreover, xv6 always keep above schema to prevent deadlock.

You should consider interrupt too.

- Sometimes, interrupt handling can cause deadlock

```
138 void
139 iderw(struct buf *b)
140 {
141     struct buf **pp;
...
150 → acquire(&idelock);
152     // Append b to idequeue.
158     // Start disk if necessary.
162     // Wait for request to finish.
163     while((b->flags &
(B_VALID|B_DIRTY)) != B_VALID){
164         sleep(b, &idelock);
165     }
166
167     release(&idelock);
```

1. Acquire
idelock

iderw function in ide.c

```
104 void
105 ideintr(void)
106 {
107     struct buf *b;
108
110     acquire(&idelock);
...
118     // Read data if needed.
122     // Wake process waiting for this
buf.
127     // Start disk on next buf in
queue.
131     release(&idelock);
```

2. IDE Interrupt is
occurred and trap
calls ideintr

3. Try to acquire
idelock but ...

ideintr function in ide.c

Solve it simple! Then is it working at last?

- So, xv6 disable interrupt during lock handling

No.

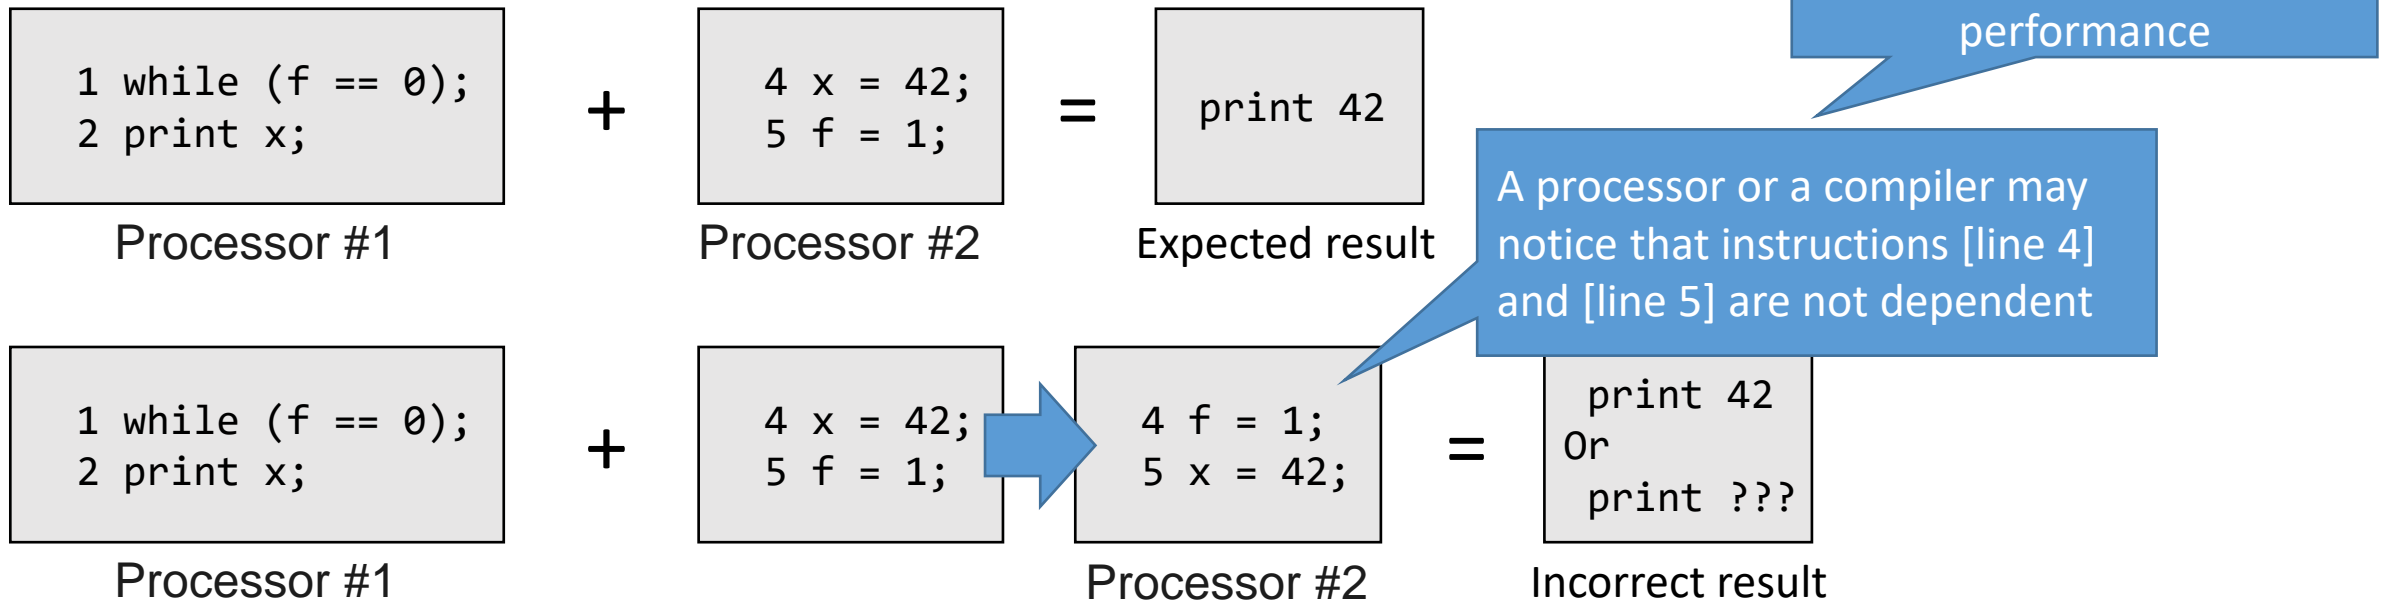
Sadly, we have one more issue

```
24 void
25 acquire(struct spinlock *lk)
26 {
27     pushcli(); // disable interrupts to avoid deadlock.
28     if(holding(lk))
29         panic("acquire");
30
31     // The xchg is atomic.
32     while(xchg(&lk->locked, 1) != 0)
33         ;
34
35     // Record info about lock acquisition for debugging.
36     lk->cpu = cpu;
37     getcallerpcs(&lk, lk->pcs);
38 }
```

```
46 void
47 release(struct spinlock *lk)
48 {
49     if(!holding(lk))
50         panic("release");
51
52     lk->pcs[0] = 0;
53     lk->cpu = 0;
54
55     // equivalent to lk->locked = 0.
56     // This code can't use a C assignment
57     // since it might not be atomic.
58     asm volatile("movl $0, %0" : "+m"
59 (lk->locked) : );
60     popcli();
61 }
```


Compiler and CPU sometimes trick you

- Many compilers and processors, however, execute code out of order to achieve higher performance.



Fortunately, We can easily deal with it

- To tell the hardware and compiler not to perform such reorderings, xv6 uses **__sync_synchronize()** in both acquire and release.
- **_sync_synchronize()** is a memory barrier
 - it tells the compiler and CPU to not reorder loads or stores across the barrier



We finally implement lock!

```
24 void
25 acquire(struct spinlock *lk)
26 {
27     pushcli(); // disable interrupts to avoid deadlock.
28     if(holding(lk))
29         panic("acquire");
30
31     // The xchg is atomic.
32     while(xchg(&lk->locked, 1) != 0)
33         ;
34     __sync_synchronize();
35
36     // Record info about lock acquisition for debugging.
37     lk->cpu = cpu;
38     getcallerpcs(&lk, lk->pcs);
39 }
```

```
46 void
47 release(struct spinlock *lk)
48 {
49     if(!holding(lk))
50         panic("release");
51
52     lk->pcs[0] = 0;
53     lk->cpu = 0;
54     __sync_synchronize();
55
56     // equivalent to lk->locked = 0.
57     // This code can't use a C assignment
58     // since it might not be atomic.
59     asm volatile("movl $0, %0" : "+m"
60 (lk->locked) : );
61     popcli();
62 }
```


Usual example of locking in xv6

```
141 fork(void)
142 {
143     int i, pid;
144     struct proc *np;
145
146     // Allocate process.
147     if((np = allocproc()) == 0){
148         return -1;
149     }
150
151     ...
174     acquire(&ptable.lock);
175
176     np->state = RUNNABLE;
177
178     release(&ptable.lock);
179
180     return pid;
181 }
```

Unusual example of locking in xv6

```
280 scheduler(void)
281 {
284     for(;;){
285         // Enable interrupts on this processor.
286         sti();
289         acquire(&ptable.lock);
290         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
291             if(p->state != RUNNABLE)
292                 continue;
297         proc = p;
298         switchvm(p);
299         p->state = RUNNING;
300         swtch(&cpu->scheduler, p->context);
301         switchkvm();
302
303         // Process is done running for now.
304         // It should have changed its p->state before coming back.
305         proc = 0;
306     }
307     release(&ptable.lock);
309 }
310 }
```

Q : Switch changes context.
So who release the lock?

A : Next process which
occupies CPU

Ptable.lock is released in ...

A fork child's very first scheduling

```
349 void
350 forkret(void)
351 {
352     static int first = 1;
353     // Still holding ptable.lock from scheduler.
354     release(&ptable.lock);
...
365     // Return to "caller", actually trapret (see allocproc).
366 }
```

Not first time

```
338 void
339 yield(void)
340 {
341     acquire(&ptable.lock);
//DOC: yieldlock
342     proc->state = RUNNABLE;
343     sched();
344     release(&ptable.lock);
345 }
```

In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281  acquire(&ptable.lock);
282  for(;;){
...
285    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292      kfree(p->kstack);
293      p->kstack = 0;
...
300      release(&ptable.lock);
301      return pid;
302    }
303  }
...
```

In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281 acquire(&ptable.lock);
282 for(;;){
...
285   for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292     kfree(p->kstack);
293     p->kstack = 0;
...
300     release(&ptable.lock);
301     return pid;
302   }
303 }
...
```

If the number of cpus is 1,
no race condition exists?

In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281 acquire(&ptable.lock);
282 for(;;){
...
285   for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292     kfree(p->kstack);
293     p->kstack = 0;
...
300     release(&ptable.lock);
301     return pid;
302   }
303 }
...
```


If the number of cpus is 1,
no race condition exists?

NO

In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281 acquire(&ptable.lock);
282 for(;;){
...
285 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292 kfree(p->kstack);
293 p->kstack = 0;
...
300 release(&ptable.lock);
301 return pid;
302 }
303 }
...
```


process 1



In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281 acquire(&ptable.lock);
282 for(;;){
...
285 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292 kfree(p->kstack);
293 p->kstack = 0;
...
300 release(&ptable.lock);
301 return pid;
302 }
303 }
...
```

process 1




Time Interrupt occurs!

In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281 acquire(&ptable.lock);
282 for(;;){
...
285 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292 kfree(p->kstack);
293 p->kstack = 0;
...
300 release(&ptable.lock);
301 return pid;
302 }
303 }
...
```

process 2

process 1




In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281 acquire(&ptable.lock);
282 for(;;){
...
285 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292 kfree(p->kstack);
293 p->kstack = 0;
...
300 release(&ptable.lock);
301 return pid;
302 }
303 }
...
```

process 2

process 1



Race condition occurs!

In uniprocessor system ..

```
274 int
275 wait(void)
276 {
...
281  acquire(&ptable.lock);
282  for(;;){
...
285    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
...
292      kfree(p->kstack);
293      p->kstack = 0;
...
300      release(&ptable.lock);
301      return pid;
302    }
303  }
...
```

If the number of cpus is 1,
a race condition can occur.

It must be protected!

Thank You
