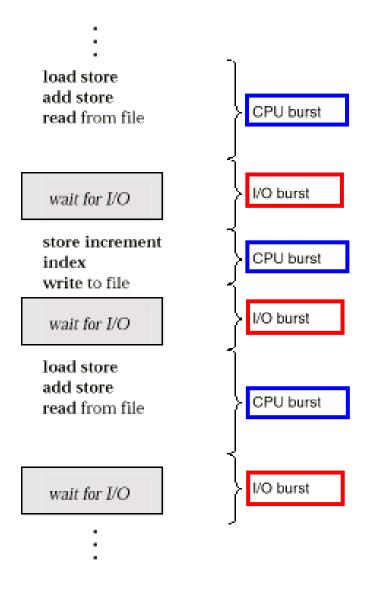


Dept. of Computer Science Hanyang University

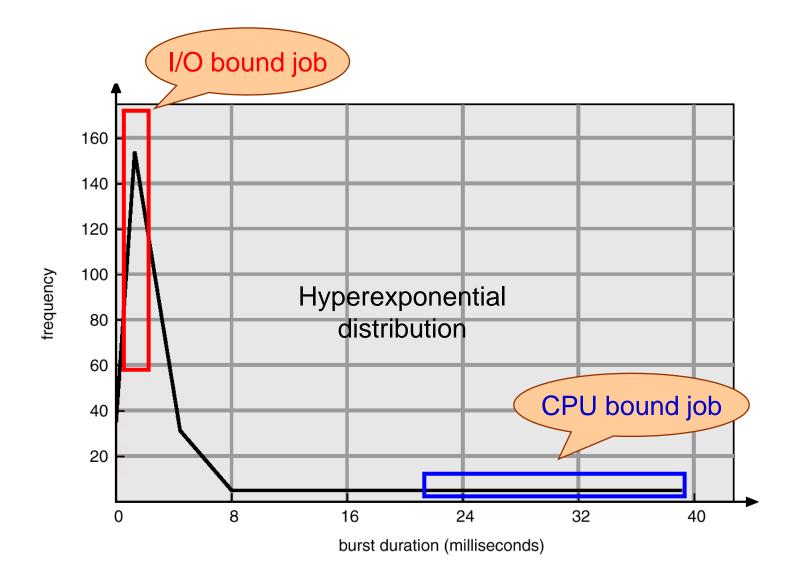




Alternating sequence of CPU and I/O Bursts



Histogram of CPU-burst Times



State Transition Diagram - revisited

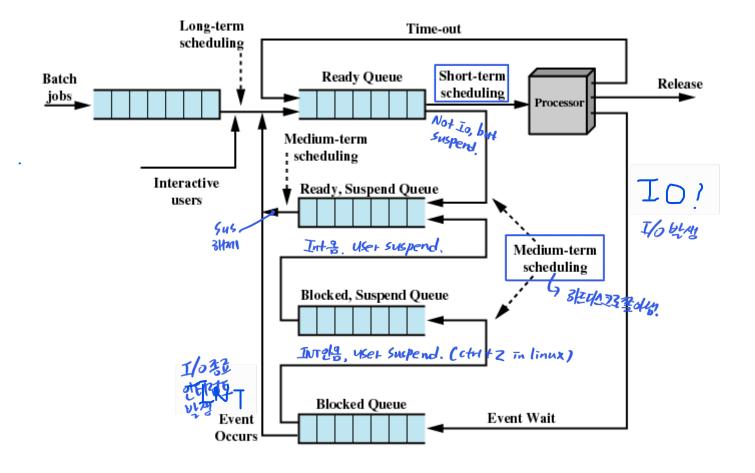


Figure 9.3 Queuing Diagram for Scheduling



Wort -> ready

어떤 placess는 court જ plan 한 것이 나?

- Multiprogramming environment
- <u>CPU Scheduler</u> selects processes in memory
 that are <u>ready</u> to execute, and <u>allocates</u> the <u>CPU</u> to one of them
- CPU scheduling decisions may take place when a process:
- Switches from running to waiting state. (eg I/O request)

2. Switches from <u>running to ready</u> state. (eg timerunout)

3. Switches from waiting to ready. (eg I/O finished interrupt)

4. Terminates

(a) Aug at court \(\frac{1}{2} \)?

(INTHANDED AT HANDE OF THE EXAMPLE THE

- If scheduling takes place only under 1 and 4:
 - The scheduling scheme is nonpreemptive
- Otherwise, preemptive scheduling scheme



 <u>Dispatcher module</u> gives control of the CPU to the process selected by the short-term scheduler, this involves:

- ✓ SWITCHING CONTEXT Context GWETCHING / TANKED APRILATING (OVERHEAD)
 - switching to user mode mak custching.
 - jumping to the proper location in the user program to restart that program

Dispatch latency

 time it takes for the dispatcher to stop one process and start another running. (Mostly, context switch overhead)



- <u>CPU utilization</u>: maximize
 - keep the CPU as busy as possible
- <u>Throughput</u>: maximize
 - # of processes that complete their execution per time unit
- Turnaround time : minimize
 - amount of time to complete a particular process
- Waiting time : minimize
 - sum of the periods spent waiting in the ready queue
- Response time: minimize Request (user Action) & tesponse-up
 - amount of time it takes from when a request was submitted <u>until the</u>
 <u>first response is produced</u>, **not** output (for time-sharing environment)



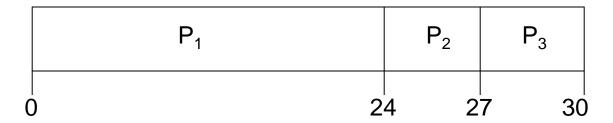
- FCFS (First-Come First-Served)
- SJF (Shortest-Job-First)
- SRTF (Shortest-Remaining-Time-First)
- Priority Scheduling
- RR (Round Robin)
- Multilevel Queue
- Multilevel Feedback Queue



FIFO의 3점 ; Fairness (39814) best/

Example:	<u>Process</u>	Burst Time
क्षेत्रस्य क्षेत्रस्	P_1	24
	P_2	3
	P_3	3

• Suppose that the processes arrive in the order: P_1 , P_2 , P_3 The Gantt Chart for the schedule is (FCFS is **nonpreemptive**):



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: (0 + 24 + 27)/3 = 17



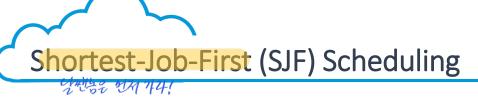
Suppose that the processes arrive in the order

$$P_2$$
, P_3 , P_1

The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: (6 + 0 + 3)/3 = 3
- Much better than previous case. watting time of
- <u>Convoy effect</u>: short process behind long process

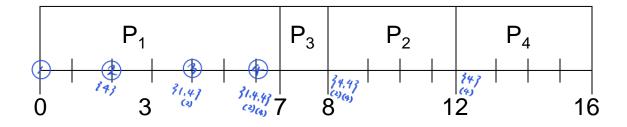


- Associate with each process the length of its next CPU burst
 - · Use these lengths to schedule the process with the shortest time
- Two schemes:
 - Nonpreemptive 时 能 to out the
 - Once CPU given to the process it cannot be preempted until completes its CPU burst
 - <u>Preemptive</u> 可器台外对? 2월 비과
 - If a new process arrives with CPU burst length less than remaining time of current executing process, preempt
 - This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal
 - Gives *minimum average waiting time* for a given set of processes



Process	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

• SJF (non-preemptive)

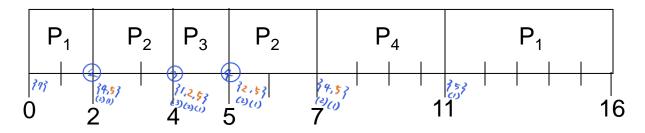


• Average waiting time = (0 + 6 + 3 + 7)/4 = 4



Process	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

• SJF (preemptive): SRTF "most optimal in waiting time"



• Average waiting time = (9 + 1 + 0 + 2)/4 = 3

$$(0+9) + (0+1) + (0) + (0+2)$$
 P_1
 P_2
 P_3
 P_a

Determining Length of Next CPU Burst

- But how do you know the length of next CPU burst?
 - ্রাচাধ মুখুর পুনাতুঃ (input data, branch, user ...)
- Can only estimate the length
- Can be done by using the length of previous CPU bursts, using exponential averaging.

```
1. t_n = \text{actuallenght of } n^{th} \text{CPU burst}
2. \tau_{n+1} = \text{predicted value for the next CPU burst}
3. \alpha, 0 \le \alpha \le 1
4. Define: \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next CPU burst \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predicted value for the next \tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n

Appendix to the predict
```

Examples of Exponential Averaging

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count.
- $\alpha = 1$
 - $\tau_{n+1} = t_n$
 - Only the actual last CPU burst counts.
- If we expand the formula, we get:

$$\begin{aligned} \tau_{n+1} &= \alpha \ t_n + (1-\alpha) \alpha \ t_{n-1} + \dots \\ &+ (1-\alpha)^{j} \alpha \ t_{n-j} + \dots \end{aligned} \qquad \text{ The first substitutions } \\ &+ (1-\alpha)^{n+1} \tau_0 = t_0 \end{aligned}$$

• Since both α and (1 - α) are less than or equal to 1, each successive term has less weight than its predecessor

Examples of Exponential Averaging

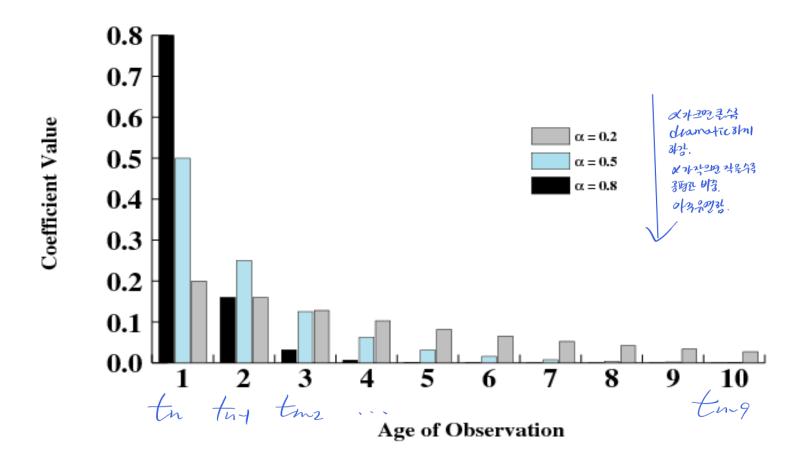
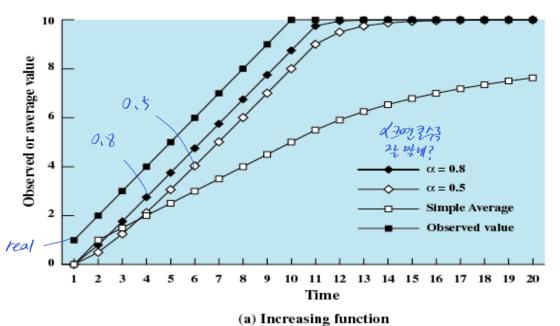


Figure 9.8 Exponential Smoothing Coefficients

大さかけかかかり



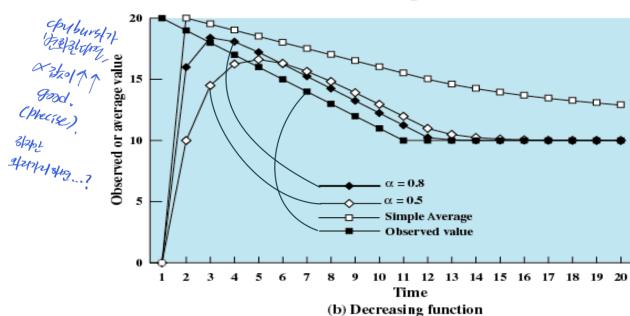


Figure 9.9 Use of Exponential Averaging



- A <u>priority number</u> (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority)
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling where
 - priority is the predicted next CPU burst time
- Problem
 - <u>Starvation</u>: low priority processes may <u>never execute</u>
- Solution
 - **Aging**: as time progresses increase the priority of the process



पपद येश्न और पीटा येमारी"

- Each process gets a small unit of CPU time (<u>time quantum</u>)
 - Usually 10-100 milliseconds
 - After this time has elapsed, the process is preempted and added to the end of the ready queue

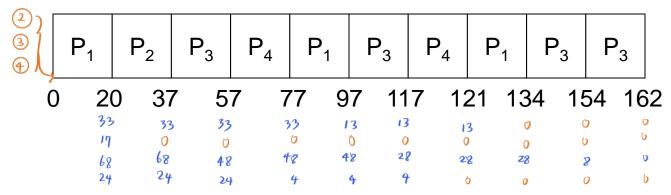
" upper bound"

- If there are *n* processes in the ready queue and the time quantum is *q sec*, then each process gets 1/*n* of the CPU time in chunks of at most *q* sec at once
 - \Rightarrow No process waits more than (n-1)q time units
- Performance
 - $q \text{ large} \Rightarrow \text{FIFO}$
 - q small ⇒ q must be large with respect to context switch,
 otherwise overhead is too high



<u>Process</u>	Burst Time
P_1	53
P_2	17
P_3	68
P_4	24

The Gantt chart is:



• Typically, higher average turnaround than SJF, but better response

- Ready queue is partitioned into separate queues:
 - <u>foreground</u> (interactive) (10 bound)
 - <u>background</u> (batch no human interaction) (cpu bound)
- Each queue has its own scheduling algorithm

```
    <u>foreground</u> – RR
    <u>background</u> – FCFS
```

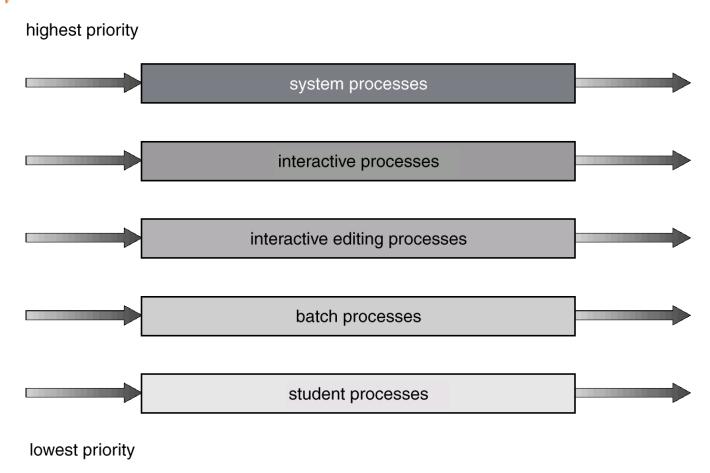
Scheduling must be done between the queues.

```
New 22 Fixed priority scheduling (interactive 2400 backsmind 31)
```

- · serve all from foreground then from background
- Possibility of <u>starvation</u>
- · Time slice the coutines foreground, we still background of the
 - each queue gets a certain amount of CPU time
 - which it can schedule amongst its processes
 - Eg., 80% to foreground in RR, 20% to background in FCFS

Multilevel Queue Scheduling

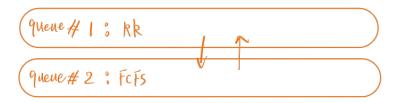
Ready queue example: queue of 574.





可吸引 E3例是子型34片 queue可写好? コ E3例とか会社queueを可容好 am 34中 protion は

- A process can move between the various queues
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service



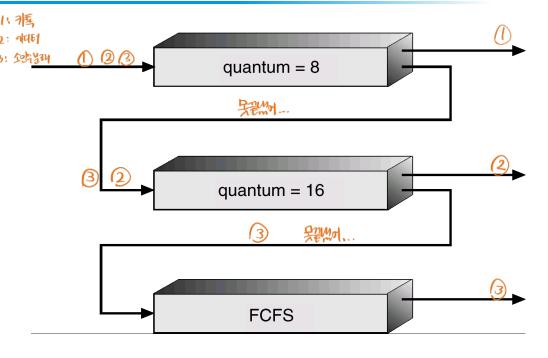
Example of Multilevel Feedback Queue

• Three queues:

- Q₀ time quantum 8 ms
- Q₁ time quantum 16 ms
- Q₂ FCFS

Scheduling

- A <u>new job enters queue Q₀</u> which is served FCFS
- When it gains CPU, job receives <u>8 milliseconds</u>
- If it does not finish in 8 milliseconds, job is moved to queue Q₁
- At Q₁ job is again served FCFS and receives <u>16 additional ms</u>
- If it still does not complete, it is preempted and moved to queue Q₂



- Hard real-time systems
 - Required to complete a critical task within a guaranteed amount of time
 - · deadline of 324. 0/2 423.
- Soft real-time computing
 - Requires that critical processes receive priority over less fortunate ones
 - ० थेममञ्च्यमान स्थाप्ता मार