




# Virtual Memory(1)

Dept. of Computer Science  
Hanyang University



# Size of a Logical Address Space




Using 32bits addresses,  $\text{max} = 2^{32} - 1$   
→ 4 GB of logical address space for each process

→ A large portion of the address space is unused



- Virtual memory: **separation** of logical memory from physical memory
  - Only part of the program needs to be in memory for execution
  - Logical address space can therefore be much larger than physical address space
  - Allows address spaces to be shared by several processes
  - Allows for more efficient process creation
  - Need to allow pages to be swapped in and out
- Virtual memory can be implemented via:
  - Demand paging
  - Demand segmentation

# Virtual Memory Larger than Physical Memory





## Demand Paging


- Bring a page into memory only when it is needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users
- Page is needed  $\Rightarrow$  reference to it
  - invalid reference  $\Rightarrow$  abort
  - not-in-memory  $\Rightarrow$  bring to memory
- Lazy swapper
  - Never swaps a page into memory unless page will be needed
  - Swapper that deals with pages is a **pager**



## Valid-Invalid Bit

- With each page table entry a valid-invalid bit is associated
  - Valid (v) : in-memory
  - "Invalid" means both:
    - *illegal*: this page is not within the address space of the process
    - *not-in-memory*: this page has never been loaded from disk before
    - *obsolete*: it is from disk, but disk copy (original) has been updated
- eg: KAL reservation system:
  - one global disk (headquater) – N computers / branch
- Initially all entries are set to *invalid*
- During address translation, if the page is `invalid'  $\Rightarrow$  **“page fault”**

# Page Table when Some Pages are not in Main Memory






# Page Fault

- Access to invalid page causes HW (MMU) trap – **page fault trap**
- Trap handler is within OS: page fault handler is invoked
- OS handles the page fault as follows:
  1. OS looks at another table to decide
    - illegal reference? eg. bad address, protection violation  $\Rightarrow$  abort process
    - not in memory? Then continue
  2. Get an empty page frame (**If no free frame, replace!**)
  3. Read the page into the frame from disk
    - The process remains in '*wait*' state until this disk I/O finishes
    - After disk I/O finishes, page table entries are updated (frame #, valid/invalid bit = "valid")
    - Move the process to the Ready queue – dispatch later
  4. Page fault trap finishes when CPU is assigned to the process, again
  5. Restart the instruction that caused the page fault



# Steps in Handling a Page Fault





## Difficulties in actual HW design

When does Page Fault occur?


1. on instruction fetch: okay

2. on operand fetch:

restart needed → (instruction fetch, decode, operand-fetch)

3. Worst case: when an instruction updates multiple locations

- Ex: Block copy instruction:



- If page fault occurs while trying to write to the second block?
  - Undo
  - Needs additional H/W that stores temporary addresses and values



## Performance of Demand Paging

- Page Fault Rate  $0 \leq p \leq 1.0$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault
- Effective Access Time (EAT)

$$\text{EAT} = (1 - p) \times \text{memory access} + p (\text{page fault overhead} + [\text{swap page out if needed}] + \text{swap page in} + \text{restart overhead})$$

- Demand paging example
  - Memory access time = 200 nanoseconds (*ns*)
  - Average page-fault service time = 8 milliseconds (*ms*)
  - $\text{EAT} = (1 - p) \times 200 + p \times 8 \text{ ms} = (1 - p) \times 200 + p \times 8,000,000 = 200 + p \times 7,999,800 \text{ ns}$
  - If one access out of 1,000 causes a page fault,  $\text{EAT} = 8.2 \mu\text{s}$
  - This is a slowdown by a factor of 40!!



## Performance of Demand Paging

- Pure demand paging
  - Never swap-in until referenced
  - Start program with no page in memory
- Locality of reference
  - Occurs in almost all workloads
  - Page references occur to very small set of pages in a certain time interval
  - Makes Paging system feasible



## What happens if there is no free frame?

- Page replacement
  - Prevent over-allocation of memory by modifying page-fault service routine to include page replacement
  - Use *modify (dirty) bit* to reduce overhead of page transfers – only modified pages are written (swap-out) to disk
  - Page replacement completes separation between logical memory and physical memory
    - Large virtual memory can be provided on a smaller physical memory
  - Same page may be brought into memory several times during run
- Page replacement algorithm
  - Algorithm for choosing victim page for replacement
  - Goal – *minimize the number of page faults*



## Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and free frame tables
4. Restart the process




## Page-Replacement Algorithms

---

- Want lowest page-fault rate
- Evaluate algorithm by running it on a particular string of memory references (**reference string**) and computing the number of page faults on that string
- In all our examples, the reference string is

**1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5**

## Graph of Page Faults v.s. the Number of Frames





## First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3** frames (3 pages can be in memory at a time per process)

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults

More page faults?

- 4** frames

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults

- FIFO Replacement – Belady’s Anomaly
  - more frames  $\Rightarrow$  more page faults



## Optimal Algorithm

- Replace page that will not be used for longest period of time
- 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	4
2	
3	
4	5

6 page faults

- How do you know this?
- Used for measuring how well your algorithm performs



## Least Recently Used (LRU) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

①	1	1	1	⑤
②	2	2	2	2
③	⑤	5	④	4
④	4	③	3	3

- Problem: How to implement LRU ?
  - If we implement LRU as it is,
    - Timestamp needed for every page: extra memory (page table) traffic
    - Need to find the page whose timestamp is the smallest
  - Too large space/time overhead to be incorporated into the Kernel
  - Approximation model for implementation needed




# LRU Implementation Algorithms

- Counter implementation
  - Every page entry has a counter;
  - CPU counter is incremented at every memory reference (logical clock)  
CPU counter = Number of total memory references
  - When page A is accessed, copy the CPU counter into the A's counter
  - At replacement, search page table for minimum counter
    - Extra memory access (counter write time) – in each memory access
    - Search (time) overhead – in each replacement
    - Counter (space) overhead, ....
- Stack implementation
  - Keep a stack of page numbers in a double link form:
  - Page A referenced:
    - move page A to the top
    - requires 6 pointers to be changed (including pointer to stack top)
  - No search for replacement


## Use of a Stack to Record the Most Recent Page References

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2



stack before a



stack after b





## LRU Approximation Algorithms

- Reference bit
  - With each page associate a bit, **Initially = 0**
  - When page is referenced, bit set to 1
  - Replace the one whose reference bit is 0 (if one exists)
    - We do not know the order, however
- Additional-Reference-Bits Algorithm
  - 8 bits for additional reference bits
  - Reference bit is shifted to the highest order bit of the additional reference bits, periodically
  - Shift the other bits right 1 bit, discarding the low-order bit
  - Ex: 00000000 – never referenced page  
10101010 – accessed in every two periods




# LRU Approximation Algorithms

- **Second chance (clock) algorithm**
  - Needs a reference bit
  - Circular queue of pages
  - Advance pointer until it finds reference bit 0 (never referenced)
  - If page to be replaced (in clock order) has reference bit = 1 then:
    - Set reference bit to 0 and leave the page in memory
    - Replace next page (in clock order), subject to same rules
  - Characteristics:
    - 포인터 이동하는 중에 reference bit 1 은 모두 0 으로 바꿈
    - 한 바퀴 되돌아와서도 (second chance) 0이면 그때에는 replace 당함
    - 자주 사용되는 페이지라면 second chance 가 올때 1
    - 최악의 경우 모든 bit 이 1 이면 FIFO 가 됨
- Enhanced Second chance algorithm

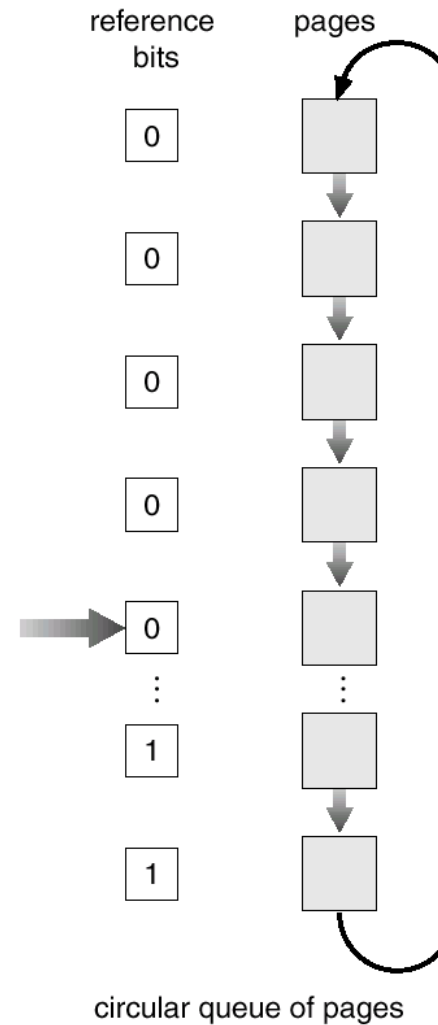
Reference bit	Modify bit	
---------------	------------	--

- |                  |              |               |
|------------------|--------------|---------------|
| • Not-Referenced | not-modified | 첫번째로 replace  |
| • Referenced     | modified     | 가장 나중 replace |

# Second-Chance (clock) Page-Replacement Algorithm



(a)



(b)





## Counting Algorithms

---

- Keep a counter of the number of references that have been made to each page
- LFU (Least Frequently Used) Algorithm
  - Replaces page with smallest count
- MFU (Most Frequently Used) Algorithm
  - Based on the argument that the page with the smallest count was probably just brought in and has yet to be used